Deep Reinforcement Learning AI

Maestracci V., Bardes A., Dezerces E.

January, 2019

Overview

Q-Learning

Deep Q-Learning

Deep Q-Learning : improvements

Our project

The principle of Q-learning

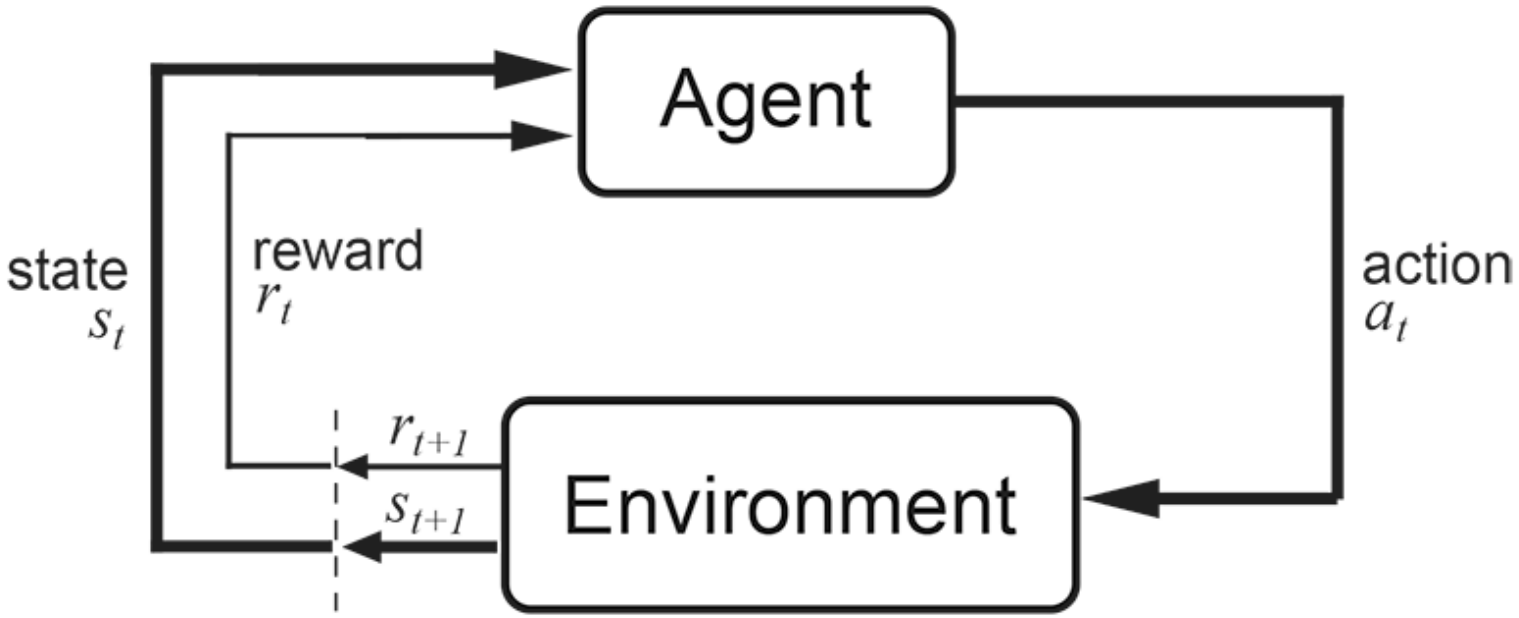

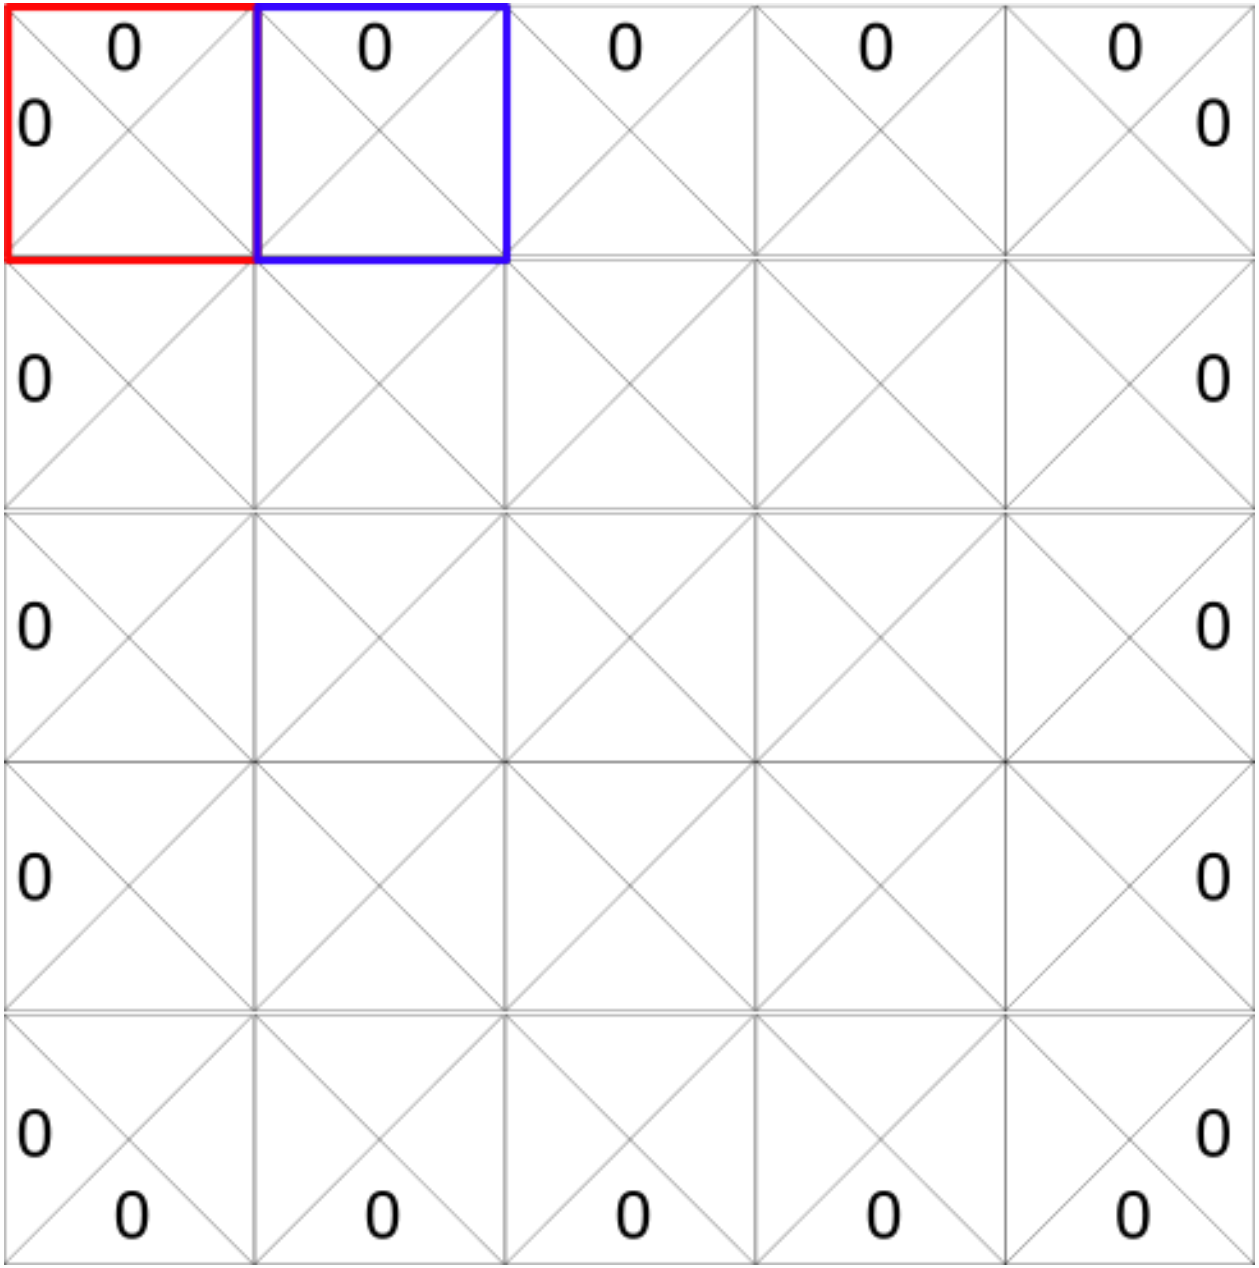
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0
0

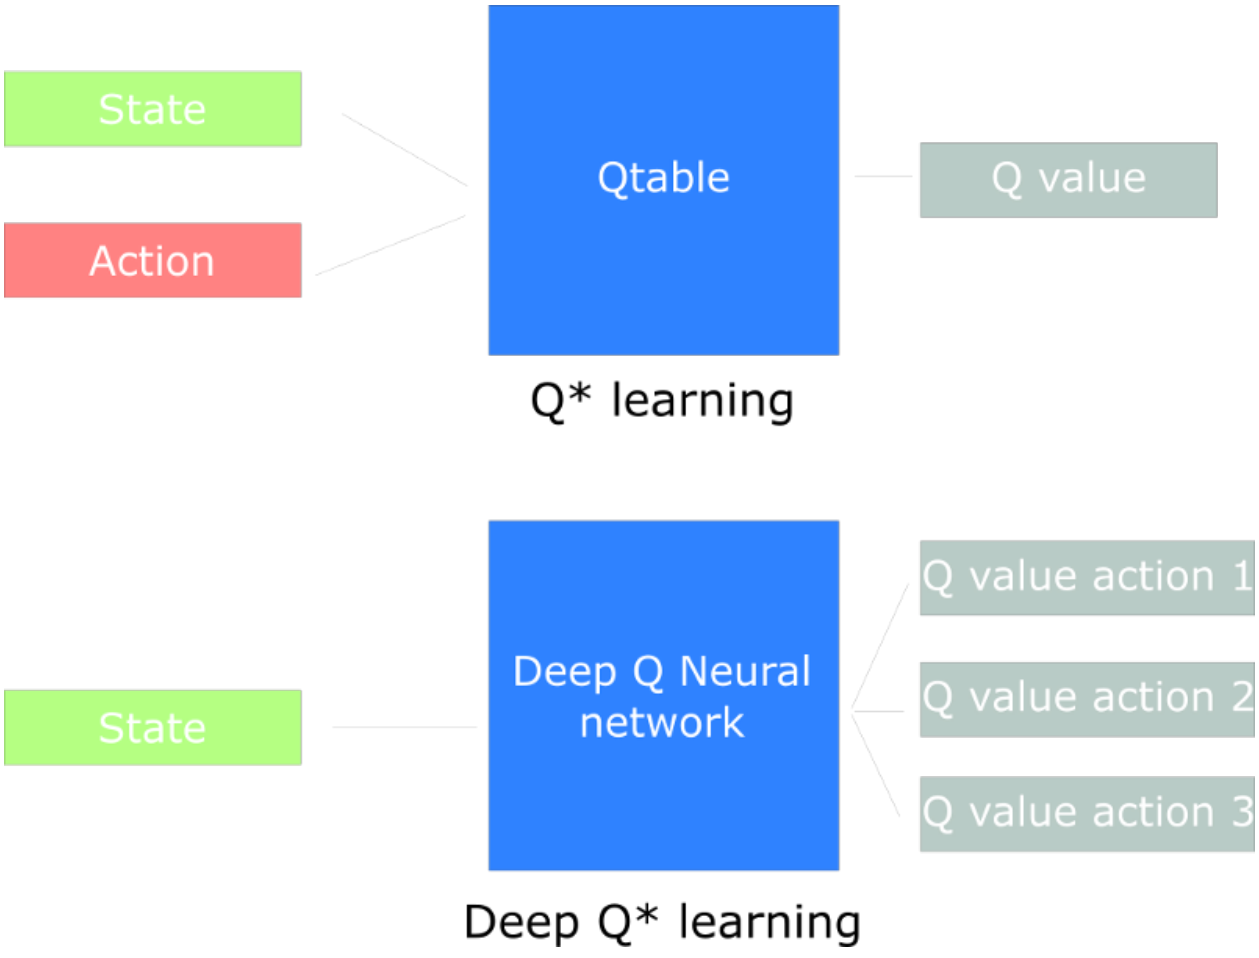
State
Action
Qtable
Q value
Q* learning
State
Deep Q Neural network
Q value action 1
Q value action 2
Q value action 3
Deep Q* learning

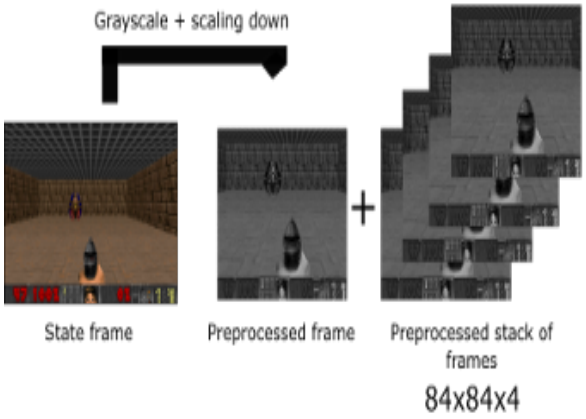
Grayscale + scaling down
State frame
Preprocessed frame
Preprocessed stack of frames
84x84x4

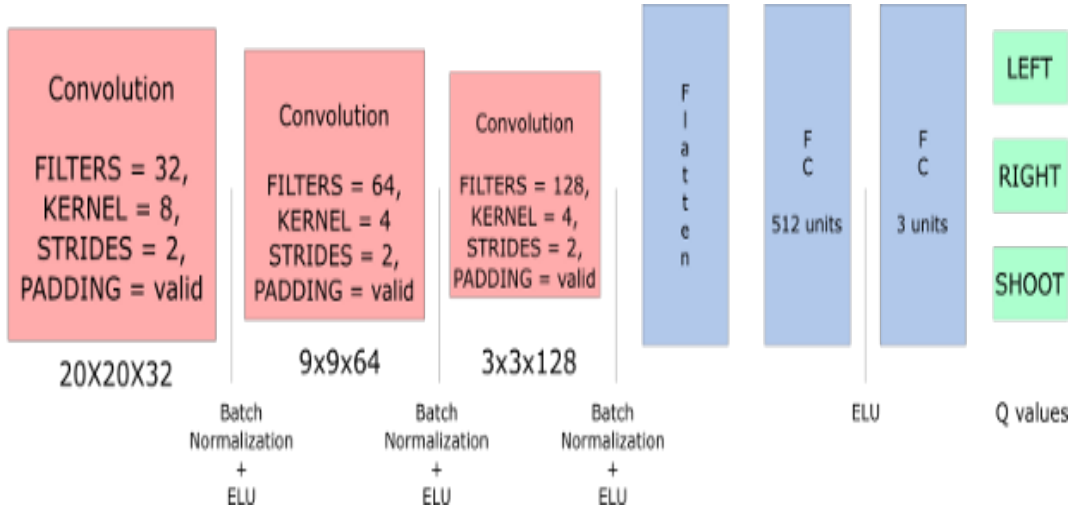
Convolution
FILTERS = 32,
KERNEL = 8,
STRIDES = 2,
PADDING = valid
20X20X32
Batch Normalization + ELU
Convolution
FILTERS = 64,
KERNEL = 4
STRIDES = 2,
PADDING = valid
9x9x64
Batch Normalization + ELU
Convolution
FILTERS = 128,
KERNEL = 4,
STRIDES = 2,
PADDING = valid
3x3x128
Batch Normalization + ELU
Flatten
FC
512 units
FC
3 units
ELU
LEFT
RIGHT
SHOOT
Q values

$$\Delta w = \alpha[(R + \gamma \, max_a \, \hat{Q}(s', a, w)) - \hat{Q}(s, a, w)] \, \nabla_w \hat{Q}(s, a, w)$$

Change in weights

learning rate

Maximum possible Qvalue for the next_state (= Q_target)

Current predicted Q-val

TD Error

Gradient of our current predicted Q-value

Memory

What if our AI goes in level two?

Memory

What if our AI goes in level two?
It musn't forget about what it learned before !

Memory

What if our AI goes in level two?
It musn't forget about what it learned before !
So let's save our previous experiences.

Fixed Q-target

We are chasing a moving target :

Fixed Q-target

We are chasing a moving target :

How can we make it so that the target doesn't move too much?

Fixed Q-target

We are chasing a moving target :

How can we make it so that the target doesn't move too much?
We can have two different networks!

Double DQN

At the beginning of the learning, our q values are noisy.
What if we favor sub-optimal choices ?

Double DQN

At the beginning of the learning, our q values are noisy.
What if we favor sub-optimal choices ?
$\longrightarrow$ We must decouple the action selection from the q-value generation.
For that we can use our two networks!

$$Q(s,a) = r(s,a) + \gamma Q(s', argmax_a Q(s',a))$$

TD target

DQN Network choose action for next state

Target network calculates the Q value of taking that action at that state

Dueling DQN

A Q value represents two things :

- How good it is to be in state s.
- How good is it to take the action a in that state.

So we can split it in two : V and A

Dueling DQN

This is the answer to the question : what is the point of knowing $Q(s, a)$ for every a when the state s is bad

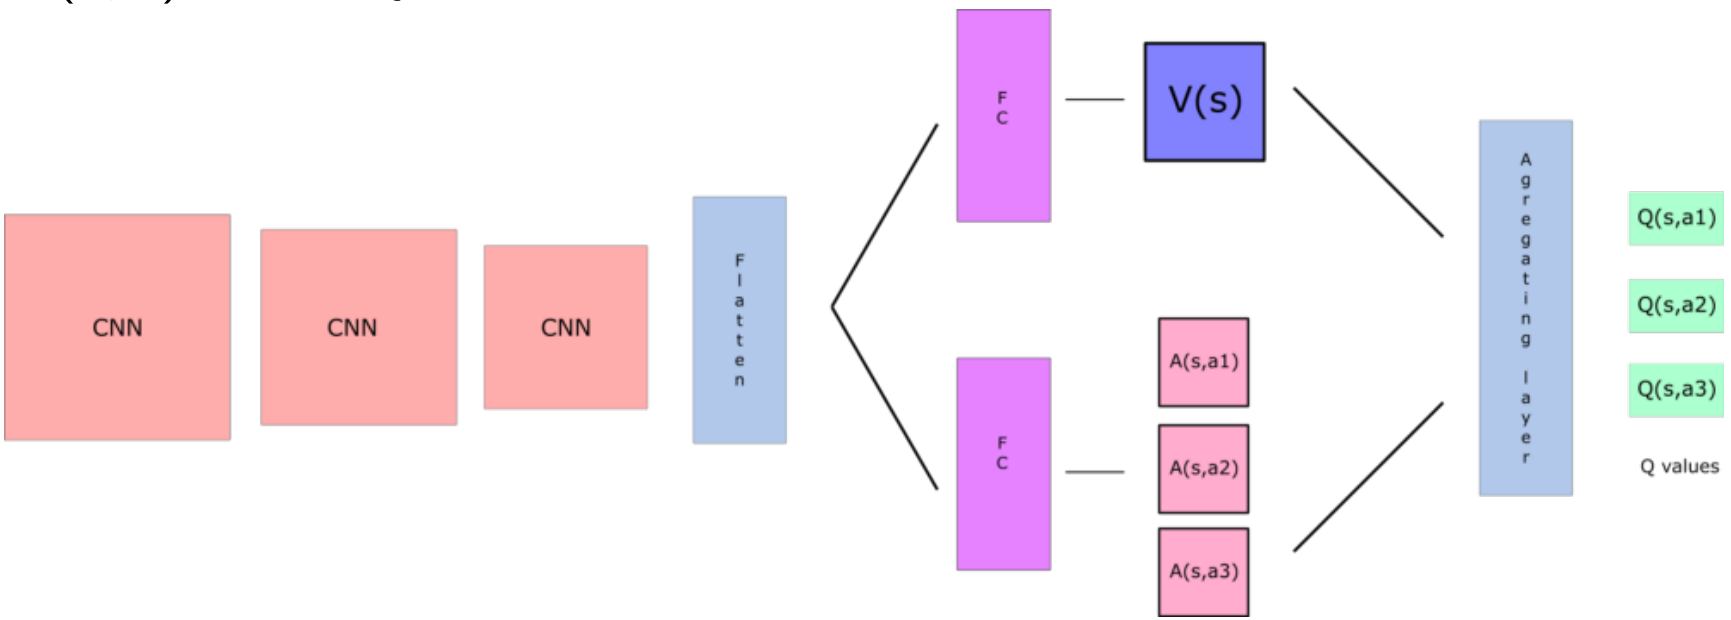

Also adds decoupling.

PER

Experiences with a huge loss are more important than others.

PER

Experiences with a huge loss are more important than others. $\longrightarrow$
So let's prioritize these.

PER

- Add a non uniform probability to be chosen for experience replay.
- Probability of being chosen decrease if you are often chosen

Unfortunately...

Our bot isn't really good...

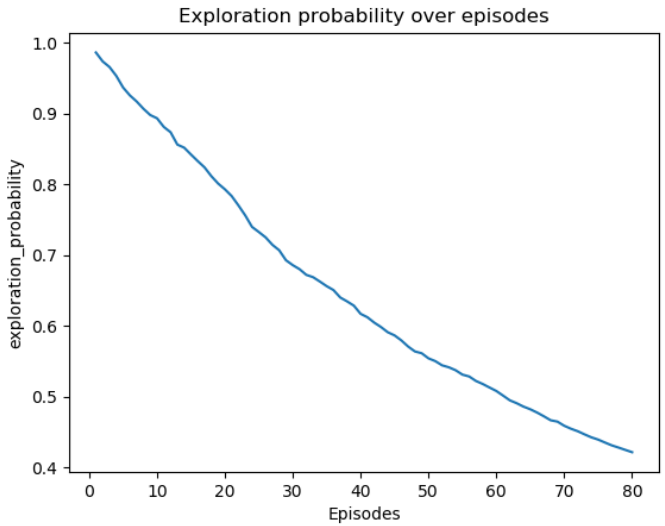
Exploration probability over episodes
exploration_probability
Episodes
1.0
0.9
0.8
0.7
0.6
0.5
0.4
0
10
20
30
40
50
60
70
80

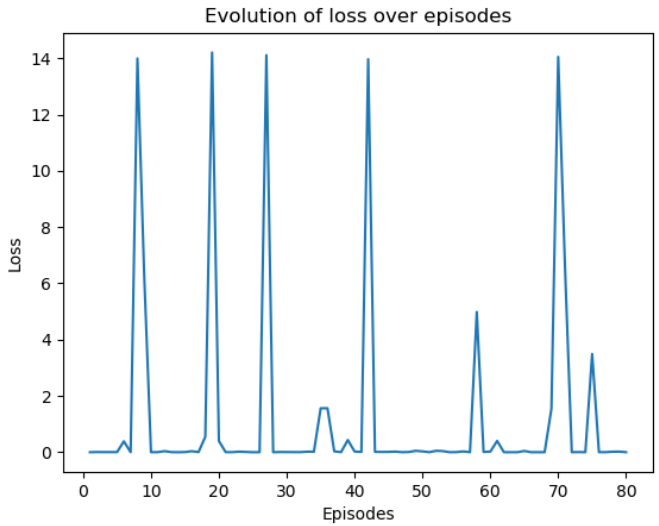
Evolution of loss over episodes
Loss
Episodes
14
12
10
8
6
4
2
0
0
10
20
30
40
50
60
70
80

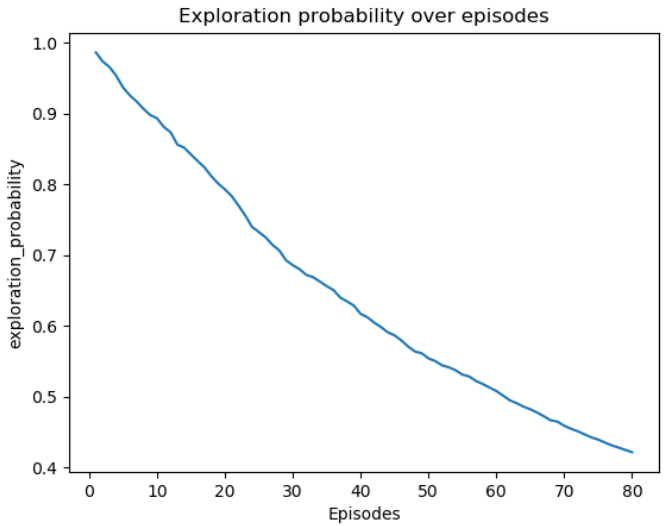
Exploration probability over episodes
exploration_probability
Episodes
1.0
0.9
0.8
0.7
0.6
0.5
0.4
0
10
20
30
40
50
60
70
80

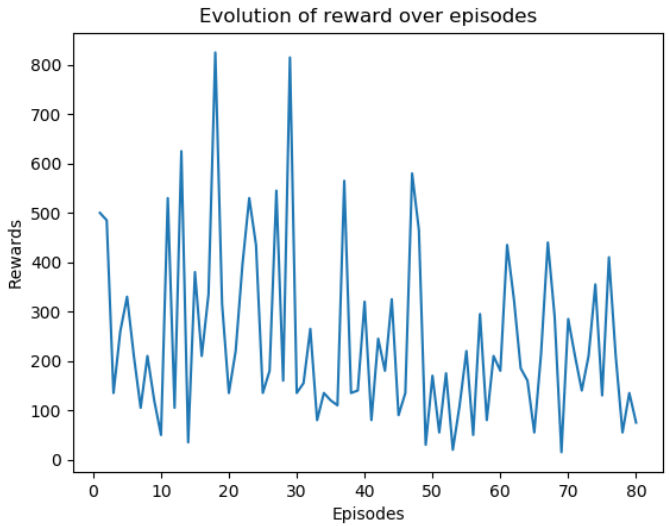
Evolution of reward over episodes
Rewards
Episodes
800
700
600
500
400
300
200
100
0
0
10
20
30
40
50
60
70
80

Our Github:
https://github.com/Adrien987k/Deep-Space-Invaders